Univerza v Ljubljani
Fakulteta za računalništvo in informatiko

Vito Levstik

Ocenjevanje volatilnosti z Markovskim prehajanjem na finančnih trgih s sekvenčnimi Monte Carlo metodami

magistrsko delo

magistrski študijski program druge stopnje
računalništvo in informatika
Smer: Podatkovne Vede

Mentor: izr. prof. dr. Jure Demšar

Ljubljana, 2026

University of Ljubljana
Faculty of Computer and Information Science

Vito Levstik

Estimating Markov switching volatility in financial markets using sequential Monte Carlo methods

master's thesis

the 2nd cycle master's study programme
computer and information science
Track: Data Science

Supervisor: izr. prof. dr. Jure Demšar

Ljubljana, 2026

This work is licensed under the *Creative Commons Attribution–ShareAlike 4.0 International (CC BY-SA 4.0)* license. This means that the text, images, graphs, and other components of this work may be freely shared, reproduced, made available to the public, and adapted for any purpose, including commercial use, provided that the author is clearly credited (preferably also the title and a link to the original), a link to the license is provided, any changes are indicated, and adaptations are distributed under the same license (CC BY-SA 4.0). The license does not permit the addition of extra legal or technological restrictions and does not apply to parts for which the rights holder is not the author. Details of the license are available at `https://creativecommons.org`.



The provided source code of the thesis, its results, and the software developed for this purpose are licensed under the GNU General Public License, version 3 (or later). This means they may be freely distributed and/or modified under its terms, and are available at `https://LINK-to-sourcecode-on-GIT`. Details of the license are available on the website at `http://www.gnu.org/licenses/`.

©2026 Vito Levstik

Acknowledgments

Worth mentioning in the acknowledgment is everyone who contributed to your thesis.

Vito Levstik, 2026

To all the flowers of this world.

"The only reason for time is so that everything doesn't happen at once."

— Albert Einstein

Contents

Abstract

Povzetek

1 Introduction **1**

2 Stochastic Volatility **3**

2.1 Volatility 3

2.2 State space models 7

2.3 Markov switching stochastic volatility model 9

2.4 Problem formulation 12

3 Sequential Monte Carlo Methods **17**

3.1 Monte Carlo method 18

3.2 Importance sampling 20

3.3 Self-normalized importance sampling 23

3.4 Sequential importance sampling 26

3.5 Sequential importance sampling with resampling 27

3.6 PMCMC 27

3.7 SMC^2 27

4 Results **29**

5 Discussion **31**

6 Conclusion **33**

A Title of the appendix 1 **35**

List of used acronyms

acronym	meaning
CA	classification accuracy
DBMS	database management system
SVM	support vector machine
...	...

Abstract

Title: Estimating Markov switching volatility in financial markets using sequential Monte Carlo methods

This sample document presents an approach to typesetting your BSc thesis using LaTeX. A proper abstract should contain around 100 words which makes this one way too short. A good abstract contains: (1) a short description of the tackled problem, (2) a short description of your approach to solving the problem, and (3) (the most successful) result or contribution in your thesis.

Keywords

computer, computer, computer

Povzetek

Naslov: Ocenjevanje volatilnosti z Markovskim prehajanjem na finančnih trgih s sekvenčnimi Monte Carlo metodami

V vzorcu je predstavljen postopek priprave magistrskega dela z uporabo okolja LaTeX. Vaš povzetek mora sicer vsebovati približno 100 besed, ta tukaj je odločno prekratek. Dober povzetek vključuje: (1) kratek opis obravnavanega problema, (2) kratek opis vašega pristopa za reševanje tega problema in (3) (najbolj uspešen) rezultat ali prispevek magistrske naloge.

Ključne besede

računalnik, računalnik, računalnik

Chapter 1

Introduction

Chapter 2

Stochastic Volatility

Volatility is a fundamental concept in financial econometrics and is commonly used as a measure of the variability or uncertainty of asset returns. In its most basic form, volatility refers to the standard deviation of returns over a given time horizon, and therefore quantifies the magnitude of fluctuations around the expected return.

Many models assume that volatility is constant. In practice, however, financial markets exhibit periods of high volatility, for example, during financial crises, and periods of low volatility, usually during stable economic conditions. Stochastic volatility models capture this behavior by allowing volatility to vary over time according to a stochastic process.

This chapter begins by introducing volatility and then develops the state space formulation used throughout the thesis, with particular emphasis on the Markov switching stochastic volatility (MSSV) model.

2.1 Volatility

Let P_t denote the price of an asset (stock, index, commodity, etc.) at time t. The log return of the asset between two consecutive time points is defined as

$$Y_t = \log(P_t) - \log(P_{t-1}) = \log\left(\frac{P_t}{P_{t-1}}\right). \tag{2.1.1}$$

The advantage of using log returns rather than simple returns lies in their additive structure over time and their natural treatment of compounding effects [1]. Thus, throughout this thesis, the term return refers to a log return.

While returns are directly observable, volatility is a characteristic of the underlying return distribution and is not directly observable from a single realization. At each time point t, only one return is observed, whereas volatility is defined through the variability of the entire return-generating distribution. Consequently, time-varying volatility must be inferred from data rather than measured directly.

To formalize this idea, we define the conditional mean and variance of the return as

$$\mu_t = \mathbb{E}[Y_t|\mathcal{F}_{t-1}], \tag{2.1.2}$$

$$\sigma_t^2 = \mathrm{Var}(Y_t|\mathcal{F}_{t-1}) = \mathbb{E}[(Y_t - \mu_t)^2|\mathcal{F}_{t-1}], \tag{2.1.3}$$

where $\mathcal{F}_{t-1}$ denotes the information set containing all relevant information available up to time $t-1$ [2]. The quantity σ_t is referred to as the conditional volatility. In many financial applications, the conditional mean is assumed to be negligible relative to the variability of returns, so that the focus is primarily on modelling the conditional volatility process.

Until now, we never specified the time scale of the returns. For example Y_t could represent daily, weekly, or monthly returns and t could therefore index days, weeks, or months. In this thesis, we focus on daily returns, specifically close-to-close daily returns, meaning that P_t is the closing price of the asset on day t[1]. The goal is then to estimate the daily conditional volatilities σ_t.

When high-frequency intraday data are available, the latent conditional variance can be approximated using realized variance. Consider a day t, divided into m equally spaced intraday intervals. Let $Y_{t,i}^{(m)}$ denote the return over the i-th intraday interval, $i = 1, \ldots, m$, so that the corresponding daily

[1] One could also consider open-to-open returns, where P_t is the opening price of the asset on day t.

return can be written as

$$Y_t = \sum_{i=1}^{m} Y_{t,i}^{(m)}. \tag{2.1.4}$$

It is often reasonable to assume that

$$\mathbb{E}\left[Y_{t,i}^{(m)} \mid \mathcal{F}_{t-1}\right] \approx 0, \tag{2.1.5}$$

$$\mathrm{Cov}\left(Y_{t,i}^{(m)}, Y_{t,j}^{(m)} \mid \mathcal{F}_{t-1}\right) \approx 0, \qquad \text{for } i \neq j. \tag{2.1.6}$$

Under these assumptions, we have

$$\sigma_t^2 = \mathrm{Var}(Y_t \mid \mathcal{F}_{t-1}) = \mathrm{Var}\left(\sum_{i=1}^{m} Y_{t,i}^{(m)} \middle| \mathcal{F}_{t-1}\right) \tag{2.1.7}$$

$$\approx \sum_{i=1}^{m} \mathrm{Var}\left(Y_{t,i}^{(m)} \mid \mathcal{F}_{t-1}\right) \tag{2.1.8}$$

$$\approx \sum_{i=1}^{m} \mathbb{E}\left[\left(Y_{t,i}^{(m)}\right)^2 \mid \mathcal{F}_{t-1}\right] \tag{2.1.9}$$

$$= \mathbb{E}\left[\sum_{i=1}^{m} \left(Y_{t,i}^{(m)}\right)^2 \mid \mathcal{F}_{t-1}\right] \tag{2.1.10}$$

$$= \mathbb{E}\left[RV_t^{(m)} \mid \mathcal{F}_{t-1}\right], \tag{2.1.11}$$

where we defined the realized variance at sampling frequency m as

$$RV_t^{(m)} = \sum_{i=1}^{m} \left(Y_{t,i}^{(m)}\right)^2. \tag{2.1.12}$$

Thus, under these assumptions, realized variance provides an approximately unbiased estimator of the conditional variance. Moreover, as the sampling frequency increases, the realized variance generally becomes a more precise measure of the underlying volatility, since the estimation error decreases with m. In practice, however, m must not be too large, to avoid that intraday returns become correlated due to market microstructure effects [3].

Many assets are not traded continuously, as markets are closed overnight, over weekends, and on holidays. We therefore decompose the daily return into overnight (close-to-open) and intraday (open-to-close) components,

$$Y_t = \log\left(\frac{P_t^{close}}{P_{t-1}^{close}}\right) = \log\left(\frac{P_t^{open}}{P_{t-1}^{close}}\right) + \log\left(\frac{P_t^{close}}{P_t^{open}}\right) = Y_t^{CO} + Y_t^{OC} \tag{2.1.13}$$

where Y_t^{CO} is the overnight (close-to-open) return and Y_t^{OC} is the intraday (open-to-close) return. Here we see the usefulness of the additive structure of log returns. From this it naturally follows that the daily variance can be decomposed into an overnight and an intraday component, if we assume that the two components are uncorrelated:

$$\sigma_t^2 = \mathrm{Var}(Y_t \mid \mathcal{F}_{t-1}) = \mathrm{Var}(Y_t^{CO} + Y_t^{OC} \mid \mathcal{F}_{t-1}) \tag{2.1.14}$$

$$= \mathrm{Var}(Y_t^{CO} \mid \mathcal{F}_{t-1}) + \mathrm{Var}(Y_t^{OC} \mid \mathcal{F}_{t-1}). \tag{2.1.15}$$

While the intraday variance can be estimated from high-frequency data using realized variance $RV_t^{(m)}$, the overnight variance is difficult to estimate reliably from a single observation[2]. One solution is to simply add the squared overnight return to $RV_t^{(m)}$. The variance estimator is then given by

$$\hat{\sigma}_t^2 = RV_t^{(m)} + (Y_t^{CO})^2. \tag{2.1.16}$$

This approach is problematic, as the squared overnight return is a very noisy estimator for the overnight variance [4].

Better approach is to scale the intraday realized variance and obtain an estimator:

$$\hat{\sigma}_t^2 = cRV_t^{(m)}, \quad c = \frac{\frac{1}{T}\sum_{t=1}^{T}(Y_t - \hat{\mu}_t)^2}{\frac{1}{T}\sum_{t=1}^{T} RV_t^{(m)}}, \tag{2.1.17}$$

where $\hat{\mu}_t$ is the sample mean of the daily returns and T is the total number of days in the sample. Here, c measures the historical ratio of total close-to-close variance to intraday variance. Equivalently, if κ denotes the historical ratio of overnight to intraday variance, then $c = 1+\kappa$. Thus, the missing overnight variance is estimated as a proportion κ of intraday realized variance rather than estimated from the noisy overnight return of each individual day. Such approach is valid under the assumption that intraday and overnight returns are uncorrelated. Weekends and holidays are naturally treated as extended overnight periods [3, 4].

[2] In case when high-frequency overnight trading data is available, the overnight variance can be approximated in the same way using intranight squared returns, just like the intraday variance [4].

Overall, when high-frequency intraday observations are available they provide a direct, increasingly precise proxy for the underlying daily variance. In empirical applications, realized variance can therefore be used in place of the unobserved conditional variance, giving us a way to evaluate the performance of volatility models.

2.2 State space models

A state space model describes a dynamical system in which an unobserved (latent) stochastic process drives the evolution of an observed process. The central idea is that the system evolves through hidden states, while only noisy measurements of those states are available.

Let $\{X_t\}_{t\geq 0}$ denote the sequence of random variables representing the latent state process, let $\{Y_t\}_{t\geq 1}$ denote the sequence of random variables representing the observed process and let T denote the time horizon, hence $t \in \{0, 1, \ldots, T\}$[3]. X_t takes values in a measurable state space $(\mathcal{X}, \mathcal{B}(\mathcal{X}))$ and Y_t takes values in $(\mathcal{Y}, \mathcal{B}(\mathcal{Y}))$. Together, $\{(X_t, Y_t)\}_{t\geq 1}$ forms a stochastic process on the joint space $(\mathcal{X} \times \mathcal{Y}, \mathcal{B}(\mathcal{X}) \otimes \mathcal{B}(\mathcal{Y}))$ [5].

The latent process $\{X_t\}_{t\geq 0}$ is modelled as a Markov process, while the observations $\{Y_t\}_{t\geq 1}$ are conditionally independent given the state process. The model is specified by the following densities:

$$X_0 \sim \chi(x_0|\theta), \tag{2.2.1}$$

$$X_t|X_{t-1} = x_{t-1} \sim p(x_t|x_{t-1}, \theta), \tag{2.2.2}$$

$$Y_t|X_t = x_t \sim f(y_t|x_t, \theta), \tag{2.2.3}$$

where χ is the *initial density*, p is the *transition density*, and f is the *observation* or *emission density*. The vector $\theta \in \Theta$ denotes the collection of model parameters [5, 6, 7].

[3]We indexed the latent states from 0 to T, whereas the observed returns are indexed from 1 to T. This way we want to emphasize that the initial state X_0 is the prior and we don't associate it with any observation. One could define the indexing differently (e.g. observations also starting with $t = 0$), but we decided to use this convention.

A potential problem with this definition is that conditional distribution of $X_t \mid X_{t-1}$ may not always admit a density. A proper way would be to define the state space model in a more measure focussed way using probability kernels [5]. However, we will not deal with this more general definition in this thesis, as our models will admit nice densities.

Sometimes, the state space model is defined in even more general terms, where the latent process and observation process are also dependent on the control input U_t, which constitutes another sequence of random variables [8]. Such definition is more used in control theory and engineering. The control inputs can be though of as external factors that influence the system dynamics and observations. In this thesis, we will not consider control inputs, as we are primarily interested in modelling the latent volatility process and its relationship with observed returns.

The main goal of state space models is to estimate the latent states based on the observed data. With $X_{i:j} = \{X_i, X_{i+1}, \ldots, X_j\}$ we denote the sequence of random variables from time i to j, and with $x_{i:j} = \{x_i, x_{i+1}, \ldots, x_j\}$ we denote their realizations. From joint posterior distribution of the latent states and model parameters given the observed data $p(x_{0:t}, \theta \mid y_{1:t})$, we can compute the following marginal distributions of interest:

- the *smoothing distribution*

$$p(x_{0:t} \mid y_{1:t}) = \int p(x_{0:t}, \theta \mid y_{1:t}) d\theta, \tag{2.2.4}$$

- the *filtering distribution*

$$p(x_t \mid y_{1:t}) = \int p(x_{0:t}, \theta \mid y_{1:t}) dx_{0:t-1} d\theta, \tag{2.2.5}$$

- the *prediction distribution*

$$p(x_{t+1} \mid y_{1:t}) = \int p(x_{t+1} \mid x_t, \theta) p(x_{0:t}, \theta \mid y_{1:t}) dx_{0:t} d\theta, \tag{2.2.6}$$

- the *posterior distribution of the parameters*

$$p(\theta \mid y_{1:t}) = \int p(x_{0:t}, \theta \mid y_{1:t}) dx_{0:t}. \tag{2.2.7}$$

A special case of state space models are the hidden Markov models (HMM). Although there is some discrepancy in the literature, the term hidden Markov model is often used to refer to a state space model in which the latent state is discrete and takes values in a finite set $\{1, 2, \ldots, K\}$, where K is the number of states [5].

2.3 Markov switching stochastic volatility model

The Markov switching stochastic volatility (MSSV) model is the main model of interest in this thesis. It is a state space model in which the hidden state is the tuple $X_t = (H_t, S_t)$. The log volatility H_t is continuous, while the regime indicator S_t is discrete and takes values in $\{1, 2, \ldots, K\}$, where K is the number of regimes [9]. The latent state space is therefore

$$\mathcal{X} = \mathbb{R} \times \{1, 2, \ldots, K\}, \tag{2.3.1}$$

equipped with the product σ-algebra $\mathcal{B}(\mathbb{R}) \otimes 2^{\{1,2,\ldots,K\}}$.

- **Observation equation:**

  The observation distribution is defined with

  $$Y_t | X_t = (h_t, s_t) \sim \mathcal{N}\left(0, \exp\left(h_t\right)\right), \tag{2.3.2}$$

  where $\mathcal{N}$ denotes the normal distribution with mean and variance [9, 10]. Such distribution admits a density function which we denote by $q(y_t|x_t, \theta) = q(y_t|h_t)$, as the observation density does not depend on the regime s_t or any other model parameters θ.

  Note that mathematically speaking the latent state H_t is the log variance of the return Y_t, but it is often referred to as the log volatility in the literature [10, 11, 12]. We will also follow this convention and refer to H_t as the log volatility.

- **Transition equation:**

  The transition distribution is defined with

$$H_t|H_{t-1} = h_{t-1}, S_t = s_t \sim \mathcal{N}(\mu_{s_t} + \phi(h_{t-1} - \mu_{s_t}), \eta^2), \tag{2.3.3}$$

$$S_t|S_{t-1} = s_{t-1} \sim \text{Categorical}([p_{s_{t-1},1}, \ldots, p_{s_{t-1},K}]), \tag{2.3.4}$$

  where μ_{s_t} denotes the regime-specific mean associated with state s_t, ϕ is the persistence parameter, η^2 is the variance of the log-volatility innovation process [9]. The values $p_{ij} = \mathbb{P}(S_t = j|S_{t-1} = i)$ are the probabilities of transitioning from regime i to regime j, and they are collected in the $K \times K$ regime transition matrix:

$$P = \begin{bmatrix} p_{11} & p_{12} & \cdots & p_{1K} \\ p_{21} & p_{22} & \cdots & p_{2K} \\ \vdots & \vdots & \ddots & \vdots \\ p_{K1} & p_{K2} & \cdots & p_{KK} \end{bmatrix},$$

  with the rows of P summing to 1, i.e., $\sum_{j=1}^{K} p_{ij} = 1$ for all i. Observe that the log volatility process is essentially an AR(1) process with regime-specific means.

  At first it seems that such conditional definition breaks the Markov property of the latent state process, as H_t depends on S_t. However, the latent state process is still Markovian, as the transition distribution of the joint state (H_t, S_t) only depends on the previous joint state (H_{t-1}, S_{t-1}). This can be seen by factoring the transition distribution as

$$p(x_t \mid x_{t-1}, \theta) = p((h_t, s_t) \mid (h_{t-1}, s_{t-1}), \theta) \tag{2.3.5}$$

$$= p(s_t \mid h_{t-1}, s_{t-1}, \theta)p(h_t \mid h_{t-1}, s_{t-1}, s_t, \theta) \tag{2.3.6}$$

$$= p(s_t \mid s_{t-1}, \theta)p(h_t \mid h_{t-1}, s_t, \theta). \tag{2.3.7}$$

- **Initial distribution:**

The MSSV model is completed by specifying the initial distribution as

$$H_0|S_0 = s_0 \sim \mathcal{N}\left(\mu_{s_0}, \frac{\eta^2}{(1-\phi^2)}\right), \tag{2.3.8}$$

$$S_0 \sim \text{Categorical}\left(\frac{1}{K}, \frac{1}{K}, \ldots, \frac{1}{K}\right), \tag{2.3.9}$$

according to the stationary distribution of the corresponding AR(1) component, with mean μ_{s_0} and variance $\frac{\eta^2}{(1-\phi^2)}$, and a uniform prior over the regimes, respectively [9].

The initial distribution admits a density function which we can write out as $\chi(x_0 \mid \theta) = \chi(s_0 \mid \theta)\chi(h_0 \mid s_0, \theta)$.

In order for the MSSV model to be well-defined, the following assumptions are made:

1. The persistence ϕ lies in the range $(-1, 1)$.
2. Each regime should correspond to at least one time point.
3. The regime means are ordered such that $\mu_1 < \mu_2 < \ldots < \mu_K$.

Assumption 1 ensures that the AR(1) process is stationary. This is needed to ensure that the log volatility does not diverge to infinity. Assumption 2 is necessary to avoid empty regimes, otherwise some μ_i would be unidentifed. Assumption 3 is a common identification constraint in mixture models, which prevents label switching and ensures that the regimes are uniquely defined [9].

In this thesis we ensure the regime identifibility using following parameterization:

$$\mu_i = \mu_{i-1} + e^{\delta_i}, \quad i = 2, \ldots, K, \tag{2.3.10}$$

where $\delta_i \in \mathbb{R}$. Such parameterization is motivated by the fact that the exponential function is strictly positive, which guarantees that $\mu_i > \mu_{i-1}$ and leaves δ_i unconstrained.

The parameters of the MSSV model are then given by vector

$$\theta = \{\mu_1, \delta_2, \ldots, \delta_K, \phi, \eta^2, P\}.$$

Since P is row stochastic, we have $K(K-1)$ free parameters in P. Therefore, the total number of free parameters is $K + 2 + K(K-1) = K^2 + 2$.

Note that in case of $K = 1$, the MSSV model reduces to the standard stochastic volatility (SV) model, in which only the log volatility H_t is a latent variable and there are no regimes.

2.4 Problem formulation

The goal of this thesis is to estimate volatility in financial markets, using the MSSV model with parameters θ. Let $y_{1:t} = \{y_1, y_2, \ldots, y_t\}$ denote the sequence of observed returns up to time $t \leq T$. Similarly, let $x_{0:t} = \{x_0, x_1, \ldots, x_t\}$ denote the sequence of latent states up to time t, where x_0 is the initial prior state. In general we can distinguish between two types of our problem.

1. **With known parameters**

   If we assume that the model parameters θ are known, then the main objective is to estimate

   $$p(x_{0:t} \mid y_{1:t}, \theta), \tag{2.4.1}$$

   which we can expand using the Bayes' theorem:

   $$p(x_{0:t} \mid y_{1:t}, \theta) \propto p(y_{1:t} \mid x_{0:t}, \theta) p(x_{0:t} \mid \theta). \tag{2.4.2}$$

We can further expand the likelihood:

$$p(y_{1:t} \mid x_{0:t}, \theta) \propto f(y_t \mid x_{0:t}, y_{1:t-1}, \theta) p(y_{1:t-1} \mid x_{0:t}, \theta) \tag{2.4.3}$$

$$\propto f(y_t \mid x_t, \theta) f(y_{t-1} \mid x_{0:t}, y_{1:t-2}, \theta) p(y_{1:t-2} \mid x_{0:t}, \theta) \tag{2.4.4}$$

$$\propto \prod_{i=1}^{t} f(y_i \mid x_i, \theta), \tag{2.4.5}$$

$$= \prod_{i=1}^{t} f(y_i \mid x_i), \tag{2.4.6}$$

where we used the fact that observations y_i depend only on the current latent state x_i (actually only on h_i) and not on the entire history of latent states, observations, or parameters. We can also expand:

$$p(x_{0:t} \mid \theta) \propto p(x_t \mid x_{0:t-1}, \theta) p(x_{0:t-1} \mid \theta) \tag{2.4.7}$$

$$\propto p(x_t \mid x_{t-1}, \theta) \dots p(x_1 \mid x_0, \theta) \chi(x_0 \mid \theta) \tag{2.4.8}$$

$$= \prod_{i=1}^{t} p(x_i \mid x_{i-1}, \theta) \chi(x_0 \mid \theta), \tag{2.4.9}$$

where we used the Markov property of the latent state process.

The final expression is thus

$$p(x_{0:t} \mid y_{1:t}, \theta) \propto \chi(x_0 \mid \theta) \prod_{i=1}^{t} p(x_i \mid x_{i-1}, \theta) f(y_i \mid x_i). \tag{2.4.10}$$

2. **With unknown parameters**

   In reality we do not know the parameters θ of our MSSV model, so we would also like to estimate them. The main objective is then the joint posterior distribution of the latent states and model parameters given the observed data:

   $$p(x_{0:t}, \theta \mid y_{1:t}), \tag{2.4.11}$$

   from which we can compute the smoothing, filtering and prediction distributions of latent states and the posterior distribution of the parameters as described in Section 2.2.

Let's expand the joint posterior distribution using the Bayes' theorem:

$$p(x_{0:t}, \theta \mid y_{1:t}) = \frac{p(x_{0:t}, y_{1:t} \mid \theta)p(\theta)}{p(y_{1:t})} \tag{2.4.12}$$

$$\propto p(y_{1:t} \mid x_{0:t}, \theta)p(x_{0:t} \mid \theta)p(\theta). \tag{2.4.13}$$

Reusing the previously derived expressions, the joint posterior distribution can then be expressed as

$$p(x_{0:t}, \theta \mid y_{1:t}) \propto p(\theta)\chi(x_0 \mid \theta) \prod_{i=1}^{t} p(x_i \mid x_{i-1}, \theta) f(y_i \mid x_i). \tag{2.4.14}$$

From both Equations (2.4.10) and (2.4.14), we can see the sequential nature of the problem. If we somehow initialy sample the prior latent state x_0 (and model parameters θ if they are unknown), we can sequentially propagate the latent states through transition density and update them with the new observations through the observation density. This is the main idea behind sequential Monte Carlo methods, which we will present in the next chapter.

2.4.1 Prior distributions of the parameters

In the case of unknown model parameters θ, we need to specify their prior distributions. We assume that the priors of the parameters are independent, so that the joint prior distribution can be factorized as

$$p(\theta) = p(\mu_1) \prod_{i=2}^{K} p(\delta_i)p(\phi)p(\eta^2)p(P). \tag{2.4.15}$$

Even though we defined our model with parameters δ_i and ϕ, we will define the priors in terms of $\Delta_i = \mu_i - \mu_{i-1} = e^{\delta_i}$ and $\varphi = \frac{\phi+1}{2}$. The reason for this reparameterization of ϕ to φ is that now φ lies in the $(0, 1)$ range, which allows us to use the Beta distribution as a prior. The reason for reparameterization of δ_i to Δ_i is that we can now use a truncated normal distribution as a prior, which will be more convenient for conditional sampling later on.

We choose the following prior distributions for the model parameters:

$$\mu_1 \sim \mathcal{N}(a_0, 5^2), \tag{2.4.16}$$

$$\Delta_i \sim \mathcal{TN}(0, \infty)(0, 5^2), \quad i = 2, \ldots, K, \tag{2.4.17}$$

$$\varphi \sim \mathcal{B}(6, 2), \tag{2.4.18}$$

$$\eta^2 \sim \mathcal{IG}(2.001, 0.1), \tag{2.4.19}$$

$$Pi, : \sim \mathcal{D}(\alpha_{i,1}, \ldots, \alpha_{i,K}), \quad i = 1, \ldots, K, \tag{2.4.20}$$

where $\alpha_{i,j} = 1.5$ for $i \neq j$ and $\alpha_{i,i} = 40$. Here $\mathcal{TN}(0, \infty)$ denotes the truncated normal distribution with support $(0, \infty)$, $\mathcal{B}$ denotes the Beta distribution, $\mathcal{IG}$ denotes the inverse gamma distribution and $\mathcal{D}$ denotes the Dirichlet distribution.

The priors are chosen to be informative while allowing sufficient flexibility. For μ_1, a_0 is determined through prior elicitation, while the variance 5^2 allows for substantial uncertainty. The prior for Δ_i allows for a wide range of possible values, but still encourages the regime means to be not be too far apart. The $\mathcal{B}(6, 2)$ prior for φ favors positive persistence, as is common in financial econometrics. Additionally, previous results from [13] suggest that the choice of prior for this parameter has little effect on posterior values, therefore we chose a value that is somewhere in between the values they considered. For η^2, we use $\mathcal{IG}(2.001, 0.1)$, following combination of values considered in [10, 13]. This prior still has a large variance, allowing for a wide range of values. For the transition matrix P, we depart from the commonly used symmetric Dirichlet prior with hyperparameters 0.5 [9, 10, 13]. Such a prior places substantial mass near 0 and 1, which may encourage regimes to either disappear or become excessively persistent. Instead, motivated by [14], we use $\alpha_{i,i} = 40$ and $\alpha_{i,j} = 1.5$ for $i \neq j$. In the two-regime case, this corresponds to a $\mathcal{B}(40, 1.5)$ prior for the diagonal transition probabilities, with mean 0.964, implying an expected regime duration of approximately 28 days. We find this expected duration reasonable for financial data.

Chapter 3

Sequential Monte Carlo Methods

Sequential Monte Carlo (SMC) methods provide a flexible framework for approximating probability distributions that evolve sequentially over time. They are particularly useful for inference in state-space models, where the latent states are typically not directly observed and their posterior distributions are analytically intractable. SMC methods approximate these distributions using a set of weighted samples, commonly referred to as particles, and update this approximation as new observations become available.

This chapter first introduces the foundations of SMC, beginning with Monte Carlo methods and sequential importance sampling. The problem of weight degeneracy is then addressed through resampling, leading to sequential importance sampling with resampling and particle filtering methods. The bootstrap particle filter and auxiliary particle filter are then presented in more detail.

The chapter then considers particle based methods for Bayesian parameter inference. Particle marginal Metropolis-Hastings and Particle Gibbs are introduced as particle Markov chain Monte Carlo methods. Finally, SMC^2 is presented as an extension that combines particle filtering with sequential inference for static parameters and latent states.

3.1 Monte Carlo method

Monte Carlo methods approximate quantities involving a probability distribution by using a finite collection of random samples. The basic idea is simple: rather than evaluating an integral directly, we replace it by an average over samples drawn from the distribution of interest.

Until further notice, let X be a random variable with probability density function $p(x)$, and let g be a measurable function such that $E_p[|g(X)|] < \infty$. The quantity of interest is

$$I = E_p[g(X)] = \int g(x)p(x)\,\mathrm{d}x. \tag{3.1.1}$$

Definition 3.1.1. Let $X_1, \ldots, X_N \overset{\text{iid}}{\sim} p(x)$. The *Monte Carlo estimator* of I is defined as

$$\hat{I}_N = \frac{1}{N}\sum_{i=1}^{N} g(X_i). \tag{3.1.2}$$

Such an estimator has several desirable properties, including unbiasedness, consistency, and asymptotic normality.

Proposition 3.1.1. *The Monte Carlo estimator $\hat{I}_N$ is unbiased.*

Proof. The proof follows from the linearity of expectation and the independence of the samples.

$$E_p[\hat{I}_N] = E_p\left[\frac{1}{N}\sum_{i=1}^{N} g(X_i)\right] \tag{3.1.3}$$

$$= \frac{1}{N}\sum_{i=1}^{N} E_p[g(X_i)] \tag{3.1.4}$$

$$= \frac{1}{N}\sum_{i=1}^{N} I = I. \tag{3.1.5}$$

□

Proposition 3.1.2. *The Monte Carlo estimator $\hat{I}_N$ is consistent.*

Proof. Because $E_p[g(X)] = I$ is finite, the strong law of large numbers applies, giving

$$\hat{I}_N \xrightarrow[N\to\infty]{a.s.} I. \tag{3.1.6}$$

Almost sure convergence implies convergence in probability, so the estimator is consistent. □

Proposition 3.1.3. *If $g(X)$ has finite variance, then the Monte Carlo estimator $\hat{I}_N$ is asymptotically normal.*

Proof. If $g(X)$ has finite variance σ^2, then the central limit theorem applies, giving

$$\sqrt{N}\left(\hat{I}_N - I\right) \xrightarrow[N\to\infty]{D} \mathcal{N}(0, \sigma^2), \tag{3.1.7}$$

which is the definition of asymptotic normality. □

3.1.1 Empirical distribution

The Monte Carlo estimator can be interpreted as integration with respect to an empirical approximation constructed from the samples.

Definition 3.1.2. Given samples $X_1, \ldots, X_N \overset{\text{iid}}{\sim} p(x)$, the *empirical distribution* associated with these samples is defined as

$$\hat{p}_N(x) = \frac{1}{N}\sum_{i=1}^{N} \delta(x - X_i), \tag{3.1.8}$$

where $\delta(\cdot)$ denotes the Dirac delta function.

The empirical distribution assigns equal weight $1/N$ to each sampled value. It therefore provides a discrete approximation to the distribution represented by $p(x)$. The expression in (3.1.8) should be understood through integration rather than as an ordinary probability density with respect to Lebesgue measure.

Proposition 3.1.4. *For any measurable function g with finite expectation,*

$$\hat{I}_N = \int g(x)\hat{p}_N(x)\,\mathrm{d}x \tag{3.1.9}$$

Proof.

$$\int g(x)\hat{p}_N(x)\,\mathrm{d}x = \int g(x)\frac{1}{N}\sum_{i=1}^{N}\delta(x - X_i)\,\mathrm{d}x \tag{3.1.10}$$

$$= \frac{1}{N}\sum_{i=1}^{N}\int g(x)\delta(x - X_i)\,\mathrm{d}x \tag{3.1.11}$$

$$= \frac{1}{N}\sum_{i=1}^{N} g(X_i) = \hat{I}_N, \tag{3.1.12}$$

where the defining property of the Dirac delta function has been used. □

It is important to distinguish the empirical distribution $\hat{p}_N$ from the density p. Even when p is continuous, $\hat{p}_N$ consists of point masses for every finite N and is therefore not a pointwise approximation of $p(x)$. The approximation is instead meaningful through quantities such as integrals and expectations.

3.2 Importance sampling

The standard Monte Carlo method requires samples from the target distribution p. In many problems, however, sampling directly from p is difficult or even impossible. Importance sampling addresses this problem by introducing a second distribution from which sampling is easier.

Definition 3.2.1. Let $q(x)$ be a proposal distribution satisfying $q(x) > 0$ whenever $p(x) > 0$. The *importance weight* associated with a point x is defined as

$$w(x) = \frac{p(x)}{q(x)}. \tag{3.2.1}$$

The expectation of g under p can then be rewritten as

$$I = E_p[g(X)] = \int g(x)p(x)\,\mathrm{d}x \tag{3.2.2}$$

$$= \int g(x)\frac{p(x)}{q(x)}q(x)\,\mathrm{d}x \tag{3.2.3}$$

$$= E_q\left[g(X)w(X)\right]. \tag{3.2.4}$$

Thus, although samples are generated from q, the importance weights correct for the difference between q and p.

Definition 3.2.2. Let $X_1, \ldots, X_N \overset{\text{iid}}{\sim} q$ and $w(X_i)$ corresponding importance weights. The *importance sampling estimator* of $I = E_p[g(X)]$ is defined by

$$\hat{I}_N^{\text{IS}} = \frac{1}{N} \sum_{i=1}^{N} g(X_i) w(X_i). \tag{3.2.5}$$

Proposition 3.2.1. *The importance sampling estimator $\hat{I}_N^{\text{IS}}$ is unbiased.*

Proof.

$$E_q[\hat{I}_N^{\text{IS}}] = \frac{1}{N} \sum_{i=1}^{N} E_q[g(X_i) w(X_i)] \tag{3.2.6}$$

$$= E_q \left[g(X) \frac{p(X)}{q(X)} \right] \tag{3.2.7}$$

$$= \int g(x) \frac{p(x)}{q(x)} q(x) \, \mathrm{d}x = I. \tag{3.2.8}$$

□

Proposition 3.2.2. *The importance sampling estimator $\hat{I}_N^{\text{IS}}$ is consistent.*

Proof. Because $E_q[g(X)w(X)] = I$ is finite, the strong law of large numbers gives

$$\hat{I}_N^{\text{IS}} \xrightarrow[N \to \infty]{a.s.} I, \tag{3.2.9}$$

which means that the estimator is consistent. □

Although the importance sampling estimator is unbiased and consistent for any proposal distribution satisfying the support condition, its variance depends strongly on the choice of q.

3.2.1 (Weighted) empirical distribution

There are two equivalent ways to interpret the importance sampling estimator. The first one follows the same logic as in 3.1.1:

$$\hat{I}_N^{\mathrm{IS}} = \frac{1}{N}\sum_{i=1}^{N} w(X_i)g(X_i) = \frac{1}{N}\sum_{i=1}^{N}\int w(x)g(x)\delta(x - X_i)\,\mathrm{d}x \tag{3.2.10}$$

$$= \int w(x)g(x)\frac{1}{N}\sum_{i=1}^{N}\delta(x - X_i)\,\mathrm{d}x \tag{3.2.11}$$

$$= \int w(x)g(x)\hat{q}_N(x)\,\mathrm{d}x, \tag{3.2.12}$$

where samples $X_1, \ldots, X_N \overset{\text{iid}}{\sim} q$ construct an empirical distribution $\hat{q}_N(x)$ and we are simply computing the expectation of $w(X)g(X)$ with respect to it.

There is another interpretation that is particularly useful. Instead of viewing the weights as part of the integrand, we can absorb them into the empirical distribution itself.

$$\hat{I}_N^{\mathrm{IS}} = \frac{1}{N}\sum_{i=1}^{N} w(X_i)g(X_i) = \frac{1}{N}\sum_{i=1}^{N}\int w(x)g(x)\delta(x - X_i)\,\mathrm{d}x \tag{3.2.13}$$

$$= \int g(x)\frac{1}{N}\sum_{i=1}^{N} w(x)\delta(x - X_i)\,\mathrm{d}x \tag{3.2.14}$$

$$= \int g(x)\hat{p}_N^{\mathrm{IS}}(x)\,\mathrm{d}x, \tag{3.2.15}$$

leading us to next definition.

Definition 3.2.3. Let $X_1, \ldots, X_N \overset{\text{iid}}{\sim} q$ and $w(X_i)$ corresponding importance weights. The *weighted empirical distribution* is defined as

$$\hat{p}_N^{\mathrm{IS}}(x) = \frac{1}{N}\sum_{i=1}^{N} w(X_i)\delta(x - X_i). \tag{3.2.16}$$

This provides an alternative interpretation of importance sampling: although the samples are drawn from the proposal distribution q, the importance weights allow the samples to represent the target distribution p. Unlike

ordinary Monte Carlo, where every sample contributes the same mass $1/N$, importance sampling assigns the sample X_i a mass proportional to $w(X_i)$.

The weighted empirical distribution in (3.2.16), however, is not generally a probability distribution for a finite sample. Its total mass is

$$\int \hat{p}_N^{\mathrm{IS}}(x)\,\mathrm{d}x = \frac{1}{N}\sum_{i=1}^{N} w(X_i), \tag{3.2.17}$$

which is a random quantity and is not necessarily equal to one. Nevertheless, the total mass has expectation equal to one. Since $X \sim q$,

$$E_q[w(X)] = \int \frac{p(x)}{q(x)} q(x)\,\mathrm{d}x = \int p(x)\,\mathrm{d}x = 1, \tag{3.2.18}$$

$$E_q\left[\frac{1}{N}\sum_{i=1}^{N} w(X_i)\right] = \frac{1}{N}\sum_{i=1}^{N} E_q[w(X_i)] = 1. \tag{3.2.19}$$

Additionally, the strong law of large numbers gives

$$\frac{1}{N}\sum_{i=1}^{N} w(X_i) \xrightarrow[N\to\infty]{a.s.} E_q[w(X)] = 1. \tag{3.2.20}$$

Thus, although the weighted empirical distribution does not necessarily have unit mass for a finite sample, its total mass converges almost surely to one as the number of samples increases.

3.3 Self-normalized importance sampling

The previous construction suggests a natural modification: instead of allowing the weights to have a random total mass, normalize them as

$$\tilde{w}_i = \frac{w(X_i)}{\sum_{j=1}^{N} w(X_j)}, \tag{3.3.1}$$

which leads us to the next definition.

Definition 3.3.1. Let $X_1, \ldots, X_N \overset{\mathrm{iid}}{\sim} q$ and $\tilde{w}(X_i)$ corresponding normalized importance weights. The *self-normalized weighted empirical distribution* is defined as

$$\hat{p}_N^{\mathrm{SNIS}}(x) = \sum_{i=1}^{N} \tilde{w}_i \delta(x - X_i) \tag{3.3.2}$$

This distribution clearly has a total mass of one.

Definition 3.3.2. Let $X_1, \ldots, X_N \overset{\text{iid}}{\sim} q$ and $w(X_i)$ corresponding importance weights. The *self-normalized imporance sampling estimator* of $I = E_p[g(X)]$ is defined as

$$\hat{I}_N^{\text{SNIS}} = \frac{\sum_{i=1}^N g(X_i)w(X_i)}{\sum_{i=1}^N w(X_i)}. \tag{3.3.3}$$

Proposition 3.3.1. *For any measurable function g with finite expectation,*

$$\hat{I}_N^{\text{SNIS}} = \int g(x)\hat{p}_N^{\text{SNIS}}\,\mathrm{d}x. \tag{3.3.4}$$

Proof.

$$\int g(x)\hat{p}_N^{\text{SNIS}}\,\mathrm{d}x = \int g(x)\sum_{i=1}^N \tilde{w}_i\delta(x - X_i)\,\mathrm{d}x \tag{3.3.5}$$

$$= \sum_{i=1}^N \tilde{w}_i \int g(x)\delta(x - X_i)\,\mathrm{d}x \tag{3.3.6}$$

$$= \sum_{i=1}^N \tilde{w}_i g(X_i) \tag{3.3.7}$$

$$= \frac{\sum_{i=1}^N g(X_i)w(X_i)}{\sum_{i=1}^N w(X_i)}. \tag{3.3.8}$$

□

There is a fundamental difference between $\hat{I}_N^{\text{IS}}$ and $\hat{I}_N^{\text{SNIS}}$. The ordinary importance sampling estimator is unbiased, whereas the self-normalized estimator is generally biased for finite N.

Proposition 3.3.2. *The self-normalized importance sampling estimator is biased.*

Proof. Because the expected value of a fraction is not equal to the fraction of expected values we have

$$E_q\left[\frac{\sum_{i=1}^N g(X_i)w(X_i)}{\sum_{i=1}^N w(X_i)}\right] \neq \frac{E_q\left[\frac{1}{N}\sum_{i=1}^N g(X_i)w(X_i)\right]}{E_q\left[\frac{1}{N}\sum_{i=1}^N w(X_i)\right]} \tag{3.3.9}$$

$$E_q\left[\hat{I}_N^{\text{SNIS}}\right] \neq \frac{E_q\left[\hat{I}_N^{\text{IS}}\right]}{1} = I. \tag{3.3.10}$$

□

Proposition 3.3.3. *The self-normalized importance sampling estimator is consistent.*

Proof. Under suitable regularity conditions, the strong law of large numbers gives

$$\frac{1}{N}\sum_{i=1}^{N} g(X_i)w(X_i) \xrightarrow[N\to\infty]{a.s.} I, \tag{3.3.11}$$

$$\frac{1}{N}\sum_{i=1}^{N} w(X_i) \xrightarrow[N\to\infty]{a.s.} 1. \tag{3.3.12}$$

Consequently,

$$\hat{I}_N^{\text{SNIS}} = \frac{\frac{1}{N}\sum_{i=1}^{N} g(X_i)w(X_i)}{\frac{1}{N}\sum_{i=1}^{N} w(X_i)} \xrightarrow[N\to\infty]{a.s.} I. \tag{3.3.13}$$

□

At first it seems pointless to use the self-normalized estimator, just to assure that the $\hat{p}_N^{\text{SNIS}}$ is a proper probability distribution, as we lose the unbiasedness of the estimator. However, in many applications, especially in Bayesian inference, the target density $p(x)$ is known only up to a proportionality constant. Suppose that our target has form

$$p(x) = \frac{\tilde{p}(x)}{Z}, \tag{3.3.14}$$

where $\tilde{p}(x)$ is known and Z is an intractable normalization constant. The importance weights are then

$$w(x) = \frac{p(x)}{q(x)} = \frac{\tilde{p}(x)}{Zq(x)}, \tag{3.3.15}$$

which makes it impossible to compute the ordinary importance sampling estimator $\hat{I}_N^{\text{IS}}$. However, the normalized weights are still computable:

$$\tilde{w}_i = \frac{w(X_i)}{\sum_{j=1}^{N} w(X_j)} = \frac{\frac{\tilde{p}(X_i)}{Zq(X_i)}}{\sum_{j=1}^{N} \frac{\tilde{p}(X_j)}{Zq(X_j)}} = \frac{\frac{\tilde{p}(X_i)}{q(X_i)}}{\sum_{j=1}^{N} \frac{\tilde{p}(X_j)}{q(X_j)}}. \tag{3.3.16}$$

Therefore, the self-normalized importance sampling estimator $\hat{I}_N^{\text{SNIS}}$ can be used even when the target density is known only up to a constant of proportionality at the cost of introducing bias.

3.4 Sequential importance sampling

Sequential importance sampling (SIS) applies importance sampling to a sequence of random variables by exploiting the sequential structure of the problem. Because the main focus of this thesis are state-space models, in particular the MSSV model, we will consider the sequential importance sampling in the context of state-space models.

Let's assume the model parameters θ are known for the moment. Given the observations $y_{1:T}$, the objective is to approximate the posterior distribution of the latent states,

$$p(x_{0:T} \mid y_{1:T}, \theta) = \frac{p(x_{0:T}, y_{1:T} \mid \theta)}{p(y_{1:T} \mid \theta)}, \tag{3.4.1}$$

where

$$p(y_{1:T} \mid \theta) = \int p(x_{0:T}, y_{1:T} \mid \theta) \, \mathrm{d}x_{0:T} \tag{3.4.2}$$

is the *marginal likelihood*. The marginal likelihood is generally intractable for nonlinear and non-Gaussian state-space models and therefore acts as an unknown normalizing constant of the posterior.

For notational simplicity, the notation for dependence on θ will be omitted. The joint density of the state-space model can be factorized according to its Markov structure as

$$p(x_{0:T}, y_{1:T}) = \chi(x_0) \prod_{t=1}^{T} f(y_t \mid x_t) p(x_t \mid x_{t-1}). \tag{3.4.3}$$

Let's assume that the proposal distribution can also be factorized sequentially as

$$q(x_{0:T} \mid y_{1:T}) = q_0(x_0) \prod_{t=1}^{T} q_t(x_t \mid x_{0:t-1}, y_{1:t}). \tag{3.4.4}$$

Then the importance weight can be computed recursively as

$$w_T = \frac{p(x_{0:T}, y_{1:T})}{q(x_{0:T} \mid y_{1:T})} = \frac{\chi(x_0) \prod_{t=1}^{T} f(y_t \mid x_t) p(x_t \mid x_{t-1})}{q_0(x_0) \prod_{t=1}^{T} q_t(x_t \mid x_{0:t-1}, y_{1:t})} \tag{3.4.5}$$

$$= w_{T-1} \cdot \frac{f(y_T \mid x_T) p(x_T \mid x_{T-1})}{q_T(x_T \mid x_{0:T-1}, y_{1:T})} = w_{T-1} \cdot \alpha_T \tag{3.4.6}$$

$$= \prod_{t=0}^{T} \alpha_t, \tag{3.4.7}$$

3.5 Sequential importance sampling with resampling

3.5.1 Bootstrap particle filter

3.5.2 Auxiliary particle filter

3.6 PMCMC

3.6.1 Particle marginal Metropolis-Hastings

3.6.2 Particle Gibbs

3.7 SMC^2

Chapter 4

Results

Chapter 5

Discussion

Chapter 6

Conclusion

Appendix A

Title of the appendix 1

Example of the appendix.

Bibliography

[1] J. Fan, Q. Yao, The Elements of Financial Econometrics, Cambridge University Press, 2017.

[2] T. G. Andersen, T. Bollerslev, P. F. Christoffersen, F. X. Diebold, Volatility and Correlation Forecasting, Vol. 1 of Handbook of Economic Forecasting, Elsevier, 2006, pp. 777–878.
URL https://www.sciencedirect.com/science/article/pii/S1574070605010153

[3] P. R. Hansen, A. Lunde, A forecast comparison of volatility models: does anything beat a GARCH(1,1)?, Journal of Applied Econometrics 20 (7) (2005) 873–889. doi:https://doi.org/10.1002/jae.800.
URL https://onlinelibrary.wiley.com/doi/abs/10.1002/jae.800

[4] M. Martens, Measuring and forecasting S&P 500 index-futures volatility using high-frequency data, Journal of Futures Markets 22 (6) (2002) 497–518. doi:https://doi.org/10.1002/fut.10016.
URL https://onlinelibrary.wiley.com/doi/abs/10.1002/fut.10016

[5] N. Chopin, O. Papaspiliopoulos, An Introduction to Sequential Monte Carlo, Springer Series in Statistics, Springer, 2020. doi:10.1007/978-3-030-47845-2.
URL https://doi.org/10.1007/978-3-030-47845-2

[6] A. Doucet, N. de Freitas, N. Gordon (Eds.), Sequential Monte Carlo Methods in Practice, Information Science and Statistics, Springer, 2001. doi:10.1007/978-1-4757-3437-9.
URL https://doi.org/10.1007/978-1-4757-3437-9

[7] V. Tingström, Sequential parameter and state learning in continuous time stochastic volatility models using the SMC^2 algorithm, Master's thesis, KTH Royal Institute of Technology (2015).
URL http://urn.kb.se/resolve?urn=urn:nbn:se:kth:diva-177104

[8] K. P. Murphy, Machine Learning: A Probabilistic Perspective, MIT Press, 2012.

[9] M. K. P. So, K. Lam, W. K. Li, A Stochastic Volatility Model with Markov Switching, Journal of Business & Economic Statistics 16 (2) (1998) 244–253.
URL http://www.jstor.org/stable/1392580

[10] C. M. Carvalho, H. F. Lopes, Simulation-based sequential analysis of Markov switching stochastic volatility models, Computational Statistics & Data Analysis 51 (9) (2007) 4526–4542.
URL https://www.sciencedirect.com/science/article/pii/S0167947306002349

[11] J. Dahlin, T. B. Schön, Getting Started with Particle Metropolis-Hastings for Inference in Nonlinear Dynamical Models, Journal of Statistical Software, Code Snippets 88 (2) (2019).
URL https://www.jstatsoft.org/index.php/jss/article/view/v088c02

[12] Y. Bao, C. Chiarella, B. Kang, Particle Filters for Markov-Switching Stochastic Volatility Models, in: The Oxford Handbook of Computational Economics and Finance, Oxford University Press, 2018.
URL https://doi.org/10.1093/oxfordhb/9780199844371.013.9

[13] Q. Pan, Y. Li, Testing volatility persistence on Markov switching stochastic volatility models, Economic Modelling 35 (2013) 45–50.
URL https://www.sciencedirect.com/science/article/pii/S0264999313002472

[14] M. Shibata, T. Watanabe, Bayesian Analysis of a Markov Switching Stochastic Volatility Model, JOURNAL OF THE JAPAN STATISTICAL SOCIETY 35 (2) (2005) 205–219. doi:10.14490/jjss.35.205.