

UNIVERZA V LJUBLJANI
FAKULTETA ZA RAČUNALNIŠTVO IN INFORMATIKO

Vito Levstik

**Ocenjevanje volatilnosti z
Markovskim prehajanjem na
finančnih trgih s sekvenčnimi Monte
Carlo metodami**

MAGISTRSKO DELO
MAGISTRSKI ŠTUDIJSKI PROGRAM DRUGE STOPNJE
RAČUNALNIŠTVO IN INFORMATIKA
SMER: PODATKOVNE VEDE

MENTOR: izr. prof. dr. Jure Demšar

Ljubljana, 2026

UNIVERSITY OF LJUBLJANA
FACULTY OF COMPUTER AND INFORMATION SCIENCE

Vito Levstik

**Estimating Markov switching
volatility in financial markets using
sequential Monte Carlo methods**

MASTER'S THESIS

THE 2ND CYCLE MASTER'S STUDY PROGRAMME
COMPUTER AND INFORMATION SCIENCE

TRACK: DATA SCIENCE

SUPERVISOR: izr. prof. dr. Jure Demšar

Ljubljana, 2026

This work is licensed under the *Creative Commons Attribution–ShareAlike 4.0 International (CC BY-SA 4.0)* license. This means that the text, images, graphs, and other components of this work may be freely shared, reproduced, made available to the public, and adapted for any purpose, including commercial use, provided that the author is clearly credited (preferably also the title and a link to the original), a link to the license is provided, any changes are indicated, and adaptations are distributed under the same license (CC BY-SA 4.0). The license does not permit the addition of extra legal or technological restrictions and does not apply to parts for which the rights holder is not the author. Details of the license are available at <https://creativecommons.org>.



The provided source code of the thesis, its results, and the software developed for this purpose are licensed under the GNU General Public License, version 3 (or later). This means they may be freely distributed and/or modified under its terms, and are available at <https://LINK-to-sourcecode-on-GIT>. Details of the license are available on the website at <http://www.gnu.org/licenses/>.

ACKNOWLEDGMENTS

Worth mentioning in the acknowledgment is everyone who contributed to your thesis.

Vito Levstik, 2026

To all the flowers of this world.

*"The only reason for time is so that
everything doesn't happen at once."*

— Albert Einstein

Contents

Abstract

Povzetek

Razširjeni povzetek	i
I Kratek pregled sorodnih del	i
II Predlagana metoda	i
III Eksperimentalna evaluacija	i
IV Sklep	i
1 Introduction	1
2 Stochastic Volatility	3
2.1 Volatility	3
2.2 State space models	5
2.3 MSSV model	5
2.4 Problem description	8
3 Sequential Monte Carlo Methods	11
3.1 Bootstrap Particle Filter	11
3.2 PMCMC	11
3.3 SMC ²	11
4 Results	13

CONTENTS

5	Discussion	15
6	Conclusion	17
A	Title of the appendix 1	19

List of used acronyms

acronym	meaning
CA	classification accuracy
DBMS	database management system
SVM	support vector machine
...	...

Abstract

Title: Estimating Markov switching volatility in financial markets using sequential Monte Carlo methods

This sample document presents an approach to typesetting your BSc thesis using L^AT_EX. A proper abstract should contain around 100 words which makes this one way too short. A good abstract contains: (1) a short description of the tackled problem, (2) a short description of your approach to solving the problem, and (3) (the most successful) result or contribution in your thesis.

Keywords

computer, computer, computer

Povzetek

Naslov: Ocenjevanje volatilnosti z Markovskim prehajanjem na finančnih trgih s sekvenčnimi Monte Carlo metodami

V vzorcu je predstavljen postopek priprave magistrskega dela z uporabo okolja L^AT_EX. Vaš povzetek mora sicer vsebovati približno 100 besed, ta tukaj je odločno prekratek. Dober povzetek vključuje: (1) kratek opis obravnavanega problema, (2) kratek opis vašega pristopa za reševanje tega problema in (3) (najbolj uspešen) rezultat ali prispevek magistrske naloge.

Ključne besede

računalnik, računalnik, računalnik

Razširjeni povzetek

To je primer razširjenega povzetka v slovenščini, ki je obvezen za naloge pisane v angleščini. Razširjeni povzetek mora vsebovati vse glavne elemente dela napisanega v angleščini skupaj s kratkim uvodom in povzetkom glavnih elementov metode, glavnih eksperimentalnih rezultatov in glavnih ugotovitev. Razširjeni povzetek naj bo strukturiran v podpoglavja (spodaj je naveden le okvirni primer in je nezavezujoč). Čez palec navadno razširjeni povzetek nanese okoli 10 odstotkov obsega celotnega dela.

I Kratek pregled sorodnih del

II Predlagana metoda

III Eksperimentalna evaluacija

IV Sklep

poljuben tekst poljuben tekst poljuben tekst poljuben tekst poljuben tekst
poljuben tekst poljuben tekst poljuben tekst poljuben tekst poljuben tekst
poljuben tekst poljuben tekst poljuben tekst poljuben tekst poljuben tekst
poljuben tekst poljuben tekst poljuben tekst poljuben tekst poljuben tekst
poljuben tekst poljuben tekst poljuben tekst poljuben tekst poljuben tekst
poljuben tekst poljuben tekst poljuben tekst poljuben tekst poljuben tekst
poljuben tekst poljuben tekst poljuben tekst poljuben tekst poljuben tekst

[illegible]

Chapter 1

Introduction

Chapter 2

Stochastic Volatility

Volatility is a fundamental concept in financial econometrics and is commonly used as a measure of the variability or uncertainty of asset returns. In its most basic form, volatility refers to the standard deviation of returns over a given time horizon, and therefore quantifies the magnitude of fluctuations around the expected return.

Many models assume that volatility is constant. In practice, however, financial markets exhibit periods of high volatility, such as during financial crises, and periods of low volatility, such as during stable economic conditions. Stochastic volatility models capture this behavior by allowing volatility to vary over time according to a stochastic process.

This chapter begins by introducing volatility and then develops the state space formulation used throughout the thesis, with particular emphasis on the Markov switching stochastic volatility (MSSV) model.

2.1 Volatility

Let P_t denote the price of an asset (stock, index, commodity, etc.) at time t . The log return of the asset is defined as

$$Y_t = \log(P_t) - \log(P_{t-1}) = \log\left(\frac{P_t}{P_{t-1}}\right). \quad (2.1)$$

The advantage of using log returns rather than simple returns lies in their additive structure over time and their natural treatment of compounding effects [1]. Throughout this thesis, the term return refers to a log return.

While returns are directly observable, volatility is a characteristic of the underlying return distribution and is not directly observable from a single realization. At each time point t , only one return Y_t is observed, whereas volatility is defined through the variability of the entire return-generating distribution. Consequently, time-varying volatility must be inferred from data rather than measured directly.

To formalize this idea, let $\mathcal{F}_{t-1}$ denote the information available up to time $t - 1$. The conditional mean and variance of the return are defined as

$$\mu_t = \mathbb{E}[Y_t | \mathcal{F}_{t-1}], \quad (2.2)$$

$$\sigma_t^2 = \text{Var}(Y_t | \mathcal{F}_{t-1}) = \mathbb{E}[(Y_t - \mu_t)^2 | \mathcal{F}_{t-1}]. \quad (2.3)$$

The quantity σ_t is referred to as the conditional volatility. In many financial applications, the conditional mean is assumed to be negligible relative to the variability of returns, so that the focus is primarily on modelling the conditional volatility process.

When high-frequency intraday data are available, volatility can be estimated nonparametrically using realized volatility. Suppose that within a fixed time interval, for example one trading day, we observe intraday log returns $r_{t,i}$ for $i = 1, \dots, m$. The realized variance is defined as

$$RV_t = \sum_{i=1}^m r_{t,i}^2, \quad (2.4)$$

Under certain regularity conditions, as the sampling frequency increases (i.e., $m \rightarrow \infty$), the realized variance converges to the quadratic variation of the log-price process, which in the absence of jumps corresponds to the integrated variance over the interval. This provides a theoretical justification for using realized volatility as a consistent estimator of latent daily volatility.

2.2 State space models

A state space model describes a dynamical system in which an unobserved (latent) stochastic process drives the evolution of an observed process. The central idea is that the system evolves through hidden states, while only noisy measurements of those states are available.

Let $\{X_t\}_{t \geq 0}$ denote the latent state process and let $\{Y_t\}_{t \geq 1}$ denote the observed process. X_t takes values in a measurable state space $(\mathcal{X}, \mathcal{B}(\mathcal{X}))$ and Y_t takes values in $(\mathcal{Y}, \mathcal{B}(\mathcal{Y}))$. Together, $\{(X_t, Y_t)\}_{t \geq 1}$ forms a stochastic process on the joint space $(\mathcal{X} \times \mathcal{Y}, \mathcal{B}(\mathcal{X}) \otimes \mathcal{B}(\mathcal{Y}))$ [2].

The latent process $\{X_t\}_{t \geq 0}$ is modelled as a Markov process, while the observations $\{Y_t\}_{t \geq 1}$ are conditionally independent given the state process. The model is specified by the following densities:

$$X_0 \sim \chi(x_0|\theta), \quad (2.5)$$

$$X_t|X_{t-1} = x_{t-1} \sim p(x_t|x_{t-1}, \theta), \quad (2.6)$$

$$Y_t|X_t = x_t \sim q(y_t|x_t, \theta), \quad (2.7)$$

where χ is the *initial density*, p is the *transition density*, and q is the *observation* or *emission density*. The vector $\theta \in \Theta$ denotes the collection of model parameters [2, 3, 4].

A potential problem with this definition is that conditional distribution of X_t given X_{t-1} may not always admit a density. A proper way would be to define the state space model in a more measure theoretic way using probability kernels [2]. However, we will not deal with this more general definition in this thesis, as our models will have nice densities.

2.3 MSSV model

The Markov switching stochastic volatility (MSSV) model is a state space model in which the hidden state is the tuple $X_t = (H_t, S_t)$. The log volatility H_t is continuous, while the regime indicator S_t is discrete and takes values in

$\{1, 2, \dots, K\}$, where K is the number of regimes [5]. The latent state space is therefore

$$\mathcal{X} = \mathbb{R} \times \{1, 2, \dots, K\}, \quad (2.8)$$

equipped with the product σ -algebra $\mathcal{B}(\mathbb{R}) \otimes 2^{\{1, 2, \dots, K\}}$.

2.3.1 Observation, transition and initial densities

The model consists of an observation equation and a transition equation. The observation distribution is defined with

$$Y_t | H_t = h_t, S_t = s_t \sim \mathcal{N}(0, \exp(h_t)), \quad (2.9)$$

where $\mathcal{N}(\mu, \sigma^2)$ denotes the normal distribution with mean μ and variance σ^2 . Such distribution admits a density function which we denote by $q(y_t | x_t, \theta) = q(y_t | h_t)$, as the observation density does not depend on the regime s_t or any other model parameters θ .

The transition distribution is defined with

$$H_t | H_{t-1} = h_{t-1}, S_t = s_t \sim \mathcal{N}(\mu_{s_t} + \phi(h_{t-1} - \mu_{s_t}), \eta^2), \quad (2.10)$$

$$S_t | S_{t-1} = s_{t-1} \sim \text{Categorical}(P_{s_{t-1}, :}), \quad (2.11)$$

where μ_{s_t} denotes the regime-specific mean associated with state s_t , ϕ is the persistence parameter, η is the standard deviation of the log-volatility innovation process, and P is the $K \times K$ regime transition matrix defined as

$$P = \begin{bmatrix} p_{11} & p_{12} & \cdots & p_{1K} \\ p_{21} & p_{22} & \cdots & p_{2K} \\ \vdots & \vdots & \ddots & \vdots \\ p_{K1} & p_{K2} & \cdots & p_{KK} \end{bmatrix},$$

where $p_{ij} = \mathbb{P}(S_t = j | S_{t-1} = i)$ is the probability of transitioning from regime i to regime j . Note that the rows of P must sum to 1, i.e., $\sum_{j=1}^K p_{ij} = 1$ for all i . Observe that the log volatility process is essentially an AR(1) process with regime-specific means [5]. Again such transition distributions admit density functions which we denote by $p(x_t | x_{t-1}, \theta) = p(s_t | s_{t-1}, \theta) p(h_t | h_{t-1}, s_t, \theta)$.

The MSSV model is completed by specifying the initial distribution. We define the initial distribution as

$$H_0|S_0 = s_0 \sim \mathcal{N}\left(\mu_{s_0}, \frac{\eta^2}{(1-\phi^2)}\right), \quad (2.12)$$

$$S_0 \sim \text{Categorical}\left(\frac{1}{K}, \frac{1}{K}, \dots, \frac{1}{K}\right). \quad (2.13)$$

The initial distribution admits a density function which we denote by $\chi(x_0|\theta) = \chi(s_0|\theta)\chi(h_0|s_0, \theta)$.

The choice of H_0 is motivated by the stationary distribution of the corresponding AR(1) component, which has mean μ_{s_0} and variance $\frac{\eta^2}{(1-\phi^2)}$. The regime variable S_0 is assigned a uniform prior because no prior information about the initial regime is assumed.

2.3.2 Assumptions and parameters

The following assumptions are made:

1. The persistence ϕ lies in the range $(-1, 1)$.
2. Each regime should correspond to at least one time point.
3. The regime means are ordered such that $\mu_1 < \mu_2 < \dots < \mu_K$.

Assumption 1 ensures that the AR(1) process is stationary. This is needed to ensure that the log volatility does not diverge to infinity. Assumption 2 is necessary to avoid empty regimes, otherwise some μ_i would be unidentified. Assumption 3 is a common identification constraint in mixture models, which prevents label switching and ensures that the regimes are uniquely defined [5].

In this thesis we ensure the regime identifiability using following parameterization:

$$\mu_i = \mu_{i-1} + e^{\delta_i}, \quad i = 2, \dots, K, \quad (2.14)$$

where $\delta_i \in \mathbb{R}$. Such parameterization is motivated by the fact that the exponential function is strictly positive, which guarantees that $\mu_i > \mu_{i-1}$ and leaves δ_i unconstrained.

The parameters of the MSSV model θ are then given by

$$\theta = \{\mu_1, \delta_2, \dots, \delta_K, \phi, \eta, P\},$$

Since P is row stochastic, we have $K(K-1)$ free parameters in P . Therefore, the total number of free parameters is $K + 2 + K(K-1) = K^2 + 2$.

Note that in case of $K = 1$, the MSSV model reduces to the standard stochastic volatility (SV) model, in which only the log volatility H_t is a latent variable and there are no regimes.

2.4 Problem description

The goal of this thesis is to estimate volatility in financial markets. Let $y_{1:t} = \{y_1, y_2, \dots, y_t\}$ denote the sequence of observed returns up to time $t \leq T$. Similarly, let $x_{0:t} = \{x_0, x_1, \dots, x_t\}$ denote the sequence of latent states up to time t , where x_0 is the initial prior state¹. The main objective is to estimate

$$p(x_{0:t}, \theta \mid y_{1:t}), \quad (2.15)$$

which is the joint posterior distribution of the latent states and model parameters given the observed returns up to time t . Integrating out the parameters θ yields the marginal posterior distribution of the latent states:

$$p(x_{0:t} \mid y_{1:t}) = \int p(x_{0:t}, \theta \mid y_{1:t}) d\theta. \quad (2.16)$$

This is known as the *smoothing distribution*. If we are interested only in the current latent state x_t , we can compute the marginal posterior distribution

¹We indexed the latent states from 0 to T , whereas the observed returns are indexed from 1 to T . This way we want to emphasize that the initial state x_0 is the prior and we don't associate it with any observed return. One could define the indexing differently (e.g. observations also starting with $t = 0$), but we decided to use this convention

of the current state:

$$p(x_t | y_{1:t}) = \int p(x_{0:t}, \theta | y_{1:t}) dx_{0:t-1} d\theta, \quad (2.17)$$

which is known as the *filtering distribution*. The filtering distribution is particularly useful for real-time applications, where we want to estimate the current state based on all available observations up to the present time. Additionally, we will also be interested in estimating the

$$p(x_{t+1} | y_{1:t}) \quad (2.18)$$

which is known as the *predictive distribution*.

Another objective is to estimate the model parameters θ from the observed data. This can be achieved by computing the marginal posterior distribution of the parameters:

$$p(\theta | y_{1:t}) = \int p(x_{0:t}, \theta | y_{1:t}) dx_{0:t}. \quad (2.19)$$

Let's expand the joint posterior distribution using Bayes' theorem:

$$p(x_{0:t}, \theta | y_{1:t}) = \frac{p(x_{0:t}, y_{1:t} | \theta) p(\theta)}{p(y_{1:t})} \quad (2.20)$$

$$\propto p(y_{1:t} | x_{0:t}, \theta) p(x_{0:t} | \theta) p(\theta). \quad (2.21)$$

We can further expand the likelihood:

$$p(y_{1:t} | x_{0:t}, \theta) \propto q(y_t | x_{0:t}, y_{1:t-1}, \theta) p(y_{1:t-1} | x_{0:t}, \theta) \quad (2.22)$$

$$\propto q(y_t | x_t, \theta) q(y_{t-1} | x_{0:t}, y_{1:t-2}, \theta) p(y_{1:t-2} | x_{0:t}, \theta) \quad (2.23)$$

$$\propto \prod_{i=1}^t q(y_i | x_i, \theta), \quad (2.24)$$

$$= \prod_{i=1}^t q(y_i | x_i), \quad (2.25)$$

where we used the fact that observations y_i depend only on the current latent state x_i (actually only on h_i) and not on the entire history of latent states,

observations, or parameters. We can also expand the factor:

$$p(x_{0:t} \mid \theta) \propto p(x_t \mid x_{0:t-1}, \theta) p(x_{0:t-1} \mid \theta) \quad (2.26)$$

$$\propto p(x_t \mid x_{t-1}, \theta) \dots p(x_1 \mid x_0, \theta) \chi(x_0 \mid \theta) \quad (2.27)$$

$$= \prod_{i=1}^t p(x_i \mid x_{i-1}, \theta) \chi(x_0 \mid \theta), \quad (2.28)$$

where we used the Markov property of the latent state process. The joint posterior distribution can then be expressed as

$$p(x_{0:t}, \theta \mid y_{1:t}) \propto p(\theta) \chi(x_0 \mid \theta) \prod_{i=1}^t p(x_i \mid x_{i-1}, \theta) q(y_i \mid x_i). \quad (2.29)$$

Chapter 3

Sequential Monte Carlo Methods

Here we present the theory of sequential Monte Carlo (SMC) methods, starting with bootstrap particle filter (BPF) and then moving to more advanced methods such as the particle Monte Carlo Markov Chain (PMCMC) and SMC².

3.1 Bootstrap Particle Filter

3.2 PMCMC

3.3 SMC²

Chapter 4

Results

Chapter 5

Discussion

Chapter 6

Conclusion

Appendix A

Title of the appendix 1

Example of the appendix.

Bibliography

- [1] J. Fan, Q. Yao, The Elements of Financial Econometrics, Cambridge University Press, 2017.
- [2] N. Chopin, O. Papaspiliopoulos, An Introduction to Sequential Monte Carlo, Springer Series in Statistics, Springer, 2020. doi:10.1007/978-3-030-47845-2.
URL <https://doi.org/10.1007/978-3-030-47845-2>
- [3] A. Doucet, N. de Freitas, N. Gordon (Eds.), Sequential Monte Carlo Methods in Practice, Information Science and Statistics, Springer, 2001. doi:10.1007/978-1-4757-3437-9.
URL <https://doi.org/10.1007/978-1-4757-3437-9>
- [4] V. Tingström, Sequential parameter and state learning in continuous time stochastic volatility models using the SMC² algorithm, Master's thesis, KTH Royal Institute of Technology (2015).
URL <http://urn.kb.se/resolve?urn=urn:nbn:se:kth:diva-177104>
- [5] M. K. P. So, K. Lam, W. K. Li, A Stochastic Volatility Model with Markov Switching, Journal of Business & Economic Statistics 16 (2) (1998) 244–253.
URL <http://www.jstor.org/stable/1392580>